

---

Workgroup: Network Working Group  
Internet-Draft: draft-epp-restful-transport-latest  
Published: 17 November 2023  
Intended: Standards Track  
Status: 20 May 2024  
Expires: M. Wullink M. Davids  
Authors: *SIDN Labs* *SIDN Labs*

# Extensible Provisioning Protocol (EPP) RESTful Transport

---

## Abstract

This document specifies a 'RESTful transport for EPP' (REPP) with the aim to improve efficiency and interoperability of EPP.

This document includes a mapping of [RFC5730] EPP commands to an HTTP based (RESTful) interface. Existing semantics and mappings as defined in [RFC5731], [RFC5732] and [RFC5733] are largely retained and reusable in RESTful EPP.

REPP allows for a stateless server implementation, no session data is stored on the EPP server. Each request from a client to the server **MUST** contain all of the information necessary for the server to process the request.

## Status of This Memo

This Internet-Draft is submitted in full conformance with the provisions of BCP 78 and BCP 79.

Internet-Drafts are working documents of the Internet Engineering Task Force (IETF). Note that other groups may also distribute working documents as Internet-Drafts. The list of current Internet-Drafts is at <https://datatracker.ietf.org/drafts/current/>.

Internet-Drafts are draft documents valid for a maximum of six months and may be updated, replaced, or obsoleted by other documents at any time. It is inappropriate to use Internet-Drafts as reference material or to cite them other than as "work in progress."

This Internet-Draft will expire on 20 May 2024.

## Copyright Notice

Copyright (c) 2023 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and

restrictions with respect to this document. Code Components extracted from this document must include Revised BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Revised BSD License.

## Table of Contents

1. Introduction	3
2. Terminology	4
3. Conventions Used in This Document	4
4. RESTful transport for EPP or REPP	5
5. Drawbacks Associated with Stateful EPP	5
6. Resource Naming Convention	5
7. Data format	6
8. Message Exchange	6
8.1. HTTP Method Definition	6
8.2. REPP Request	7
8.2.1. EPP Data	7
8.2.2. REPP Request Headers	7
8.2.3. Generic HTTP Headers	7
8.3. REPP Response	8
8.3.1. REPP Response Headers	8
8.3.2. Generic Headers	8
9. Command Mapping	9
9.1. Hello	10
9.2. Session Management Resources	11
9.2.1. Login	11
9.2.2. Logout	11
9.3. REPP Query Commands	11
9.3.1. Check	12
9.3.2. Info	13
9.3.3. Poll	13

9.4. REPP Transform Commands	14
9.4.1. Create	14
9.4.2. Delete	14
9.4.3. Renew	14
9.4.4. Update	15
9.4.5. Transfer	15
10. Transport Considerations	16
11. IANA Considerations	16
12. Internationalization Considerations	16
13. Security Considerations	16
14. Obsolete EPP Result Codes	17
15. Acknowledgments	17
16. Normative References	17
Appendix A. Examples	19
A.1. X-REPP-auth-info	19
A.1.1. Domain Info with Authorization Data	19
A.2. Domain Create Example	19
A.2.1. Domain Create Request	19
A.2.2. Domain Create Response:	21
A.3. Domain Delete Example	21
A.3.1. Domain Delete Request:	21
A.3.2. Domain Delete Response:	22
Authors' Addresses	22

## 1. Introduction

This document describes a transport protocol for EPP, based on the [\[REST\]](#) architectural style. This transport mechanism leverages the HTTP protocol [\[RFC2616\]](#) and the principles of [\[REST\]](#). Conforming to the REST constraints is generally referred to as being "RESTful". Hence we dubbed the new transport protocol: "RESTful transport for EPP" or "REPP" for short.

This new REST based transport includes a mapping of [\[RFC5730\]](#) EPP-commands to [\[URI\]](#) resources. REPP, in contrast to the EPP specification, is stateless. It aims to provide a mechanism that is more suitable for complex, high availability environments, as well as for environments where TCP connections can be unreliable.

RFC 5730 [\[RFC5730\]](#) Section 2.1 describes that EPP can be layered over multiple transport protocols. Currently, the EPP transport over TCP [\[RFC5734\]](#) is the only widely deployed transport mapping for EPP. This same section defines that newly defined transport mappings must preserve the stateful nature of EPP.

The stateless nature of REPP dictates that no session state is maintained on the EPP server. Each request from client to server will contain all of the information necessary to understand the request. The server will close the connection after each HTTP request.

With a stateless mechanism, some drawbacks of EPP (as mentioned in Section 5) are circumvented.

A good understanding of the EPP base protocol specification [\[RFC5730\]](#) is advised, to grasp the command mapping described in this document.

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [\[RFC2119\]](#).

## 2. Terminology

In this document the following terminology is used.

REST - Representational State Transfer ([\[REST\]](#)). An architectural style.

RESTful - A RESTful web service is a web service or API implemented using HTTP and the principles of [\[REST\]](#).

EPP RFCs - This is a reference to the EPP version 1.0 specifications [\[RFC5730\]](#), [\[RFC5731\]](#), [\[RFC5732\]](#) and [\[RFC5733\]](#).

Stateful EPP - The definition according to Section 2 of [\[RFC5730\]](#).

RESTful EPP or REPP - The RESTful transport for EPP described in this document.

URL - A Uniform Resource Locator as defined in [\[RFC3986\]](#).

Resource - A network data object or service that can be identified by a URL.

Command mapping - The mapping of [\[RFC5730\]](#) EPP commands to RESTful EPP.

## 3. Conventions Used in This Document

XML is case sensitive. Unless stated otherwise, XML specifications and examples provided in this document MUST be interpreted in the character case presented to develop a conforming implementation.

The examples in this document assume that request and response messages are properly formatted XML documents.

## 4. RESTful transport for EPP or REPP

REPP is designed to solve, in the spirit of [\[RFC3375\]](#), the drawbacks as mentioned in the next paragraph and yet maintain compatibility with existing object mapping definitions.

The design intent is to provide a clear, clean and self-explanatory interface that can easily be integrated with existing software systems. On the basis of these principles a [REST] architectural style was chosen. A client interacts with a REPP server via HTTP requests.

A server implementing REPP, MUST NOT keep any client state. Every client request needs to provide all of the information necessary to process the request.

REPP conforms to the EPP transport mapping considerations as defined in [\[RFC5730\]](#), Section 2.1. With REPP, the EPP [\[RFC5730\]](#) commands are mapped to REST URL resources.

## 5. Drawbacks Associated with Stateful EPP

[\[RFC5734\]](#) requires a stateful session between a client and the EPP server. This is accomplished by setting up a session with a <login> and keeping it alive for some time until issuing a <logout>. This may pose challenges in load-balanced environments, when a running session for whatever reason suddenly has to be switched from one EPP server to another and state is kept on a per server basis.

[\[RFC5734\]](#) EPP sessions can wind up in a state where they are no longer linked to an active TCP connection, especially in an environment where TCP connectivity is flaky. This may raise problems in situations where session limits are enforced.

REPP is designed to avoid these drawbacks, hence making the interaction between an EPP client and an EPP server more robust and efficient.

## 6. Resource Naming Convention

A resource can be a single uniquely object identifier e.g. a domain name, or a collection of objects. The complete set of objects available to client for registry operations MUST be identified by {context- root}/{version}/{collection}

- o {context-root} is the base URL which MUST be specified by each registry. The {context-root} MAY be an empty, zero length string.

- o {version} is a label which identifies the interface version. This is the equivalent of the <version> element in the EPP RFCs. The version used in a REPP URL MUST match the version used by EPP in the upper layer.

- o {collection} MUST be substituted by "domains", "hosts" or "contacts", referring to either [\[RFC5731\]](#), [\[RFC5732\]](#) or [\[RFC5733\]](#).

- o A trailing slash MAY be added to each request. Implementations MUST consider requests which only differ with respect to this trailing slash as identical.

A specific object instance MUST be identified by {context-root}/{version}/{collection}/{id} where {id} is a unique object identifier described in EPP RFCs.

An example domain name resource, for example.nl, would look like this:

```
/repp/v1/domains/example.nl
```

The level below a collection MUST be used to identify an object instance, the level below an object instance MUST be used to identify attributes of the object instance.

With REPP the object identifiers are embedded in URLs. This makes any object identifier in the request messages superfluous. However, since the goal of REPP is to stay compatible with the existing EPP object mapping schemas, this redundancy is accepted as a trade off. Removing the object identifier from the request message would require new object mapping schemas.

The server MUST return HTTP Status-Code 412 when the object identifier (for example [domain:name](#), [host:name](#) or [contact:id](#)) in the HTTP message-body does not match the {id} object identifier in the URL.

## 7. Data format

TODO Describe how REP MUST be data format agnostic and support multiple data formats. XML is already defined, but future JSON mappings MUST also be supported.

now using http header: Accept: application/epp+xml MUST support other content-types  
e.g.: application/epp+json

## 8. Message Exchange

A [\[RFC5730\]](#) EPP request includes a command- and object mapping to which a command must be applied. REPP uses the REST semantics each HTTP method is assigned a distinct behaviour, section [Section 8.1](#) provides a overview of each the behaviour assigned to each method. REPP requests are expressed by using an URL resource, an HTTP method and optionally an HTTP header and message body.

Data (payload) belonging to a request or response is added to the HTTP message- body or sent as using an HTTP header, depending on the nature of the request as defined in Section 9.

An HTTP request MUST contain no more than one EPP command. HTTP requests MUST be processed independently of each other and in the same order as the server receives them.

### 8.1. HTTP Method Definition

REPP commands MUST be executed by using an HTTP method on a resource identified by an URL. The server MUST support the following "verbs" ([\[REST\]](#)).

GET: Request a representation of a resource or a collection of resources.

PUT: Update an existing resource.

POST: Create a new resource.

DELETE: Delete an existing resource.

HEAD: Check for the existence of a resource.

OPTIONS: Request a greeting.

The server **MUST** not support the following "verbs"

PATCH: Partial updating of a resource is **MUST** not be allowed.

## **8.2. REPP Request**

TODO: add text that due to stateless nature a request must contain all data required to execute command. includes data previously part of client session (login)

### **8.2.1. EPP Data**

The payload data for a REPP request **MAY** be transmitted to the server using the POST, PUT and GET HTTP methods.

POST and PUT: Payload data, when required, **MUST** be added to the message-body.

GET: When payload data is required, it concerns <authInfo>. This **SHALL** be put in the "X-REPP-auth-info" HTTP request-header.

### **8.2.2. REPP Request Headers**

HTTP request-headers are used to transmit additional or optional request data to the server. All REPP HTTP headers must have the "X-REPP-" prefix.

X-REPP-cltrid: The client transaction identifier is the equivalent of the <clTRID> element in the EPP RFCs and **MUST** be used accordingly. When this header is present in a client request, an equivalent element in the message-body **MAY** also be present, but **MUST** than be consistent with the header.

X-REPP-svcs: The namespace used by the client in the EPP request document in the message-body of the HTTP request.

### **8.2.3. Generic HTTP Headers**

Generic HTTP headers **MAY** be used as defined in HTTP/1.1 [[RFC2616](#)]. For REPP, the following general-headers are **REQUIRED** in HTTP requests.

Accept-Language: This request-header is equivalent to the <lang> element in the EPP <login> command, expect that the usage of this header by the client is **OPTIONAL**. The server **MUST** support the use of HTTP Accept-Language header in client requests. The client **MAY** issue a <hello> to discover the languages known by the server. Multiple servers in a load-balanced environment **SHOULD** reply with consistent <lang> elements in a <greeting>. Clients **SHOULD NOT** expect that obtained <lang> information remains consistent between different requests. Languages not supported by the server default to "en".

Content-Type: ...

### 8.3. REPP Response

The server response contains an HTTP Status-Code, HTTP response-headers and it MAY contain an EPP response message in the HTTP message-body.

#### 8.3.1. REPP Response Headers

HTTP response-headers are used to transmit additional response data to the client. All HTTP headers used by REPP MUST use the "X-REPP-" prefix.

X-REPP-svtrid: This header is the equivalent of the <svTRID> element in the EPP RFCs and MUST be used accordingly. If an HTTP message-body with the EPP XML equivalent <svTRID> exists, both values MUST be consistent.

X-REPP-cltrid: This header is the equivalent of the <clTRID> element in the EPP RFCs and MUST be used accordingly. If an HTTP message-body with the EPP XML equivalent <clTRID> exists, both values MUST be consistent.

X-REPP-eppcode: This header is the equivalent of the <result code> element in the EPP RFCs and MUST be used accordingly. If an HTTP message-body with the EPP XML equivalent <result code> exists, both values MUST be consistent.

X-REPP-check-avail: An alternative for the "avail" attribute of the [object:name](#) element in a check response and MUST be used accordingly.

X-REPP-check-reason: An optional alternative for the "object:reason" element in a check response and MUST be used accordingly.

#### 8.3.2. Generic Headers

Generic HTTP headers MAY be used as defined in HTTP/1.1 [[RFC2616](#)]. For REPP, the following general-headers are REQUIRED in HTTP responses.

Cache-Control: ... TBD: the idea is to prohibit caching. Even though it will probably work and be useful in some scenario's, it also complicates matters.]

Connection: ....

#### ## Error Handling

REPP is designed atop of the HTTP protocol, both are an application layer protocol with their own status- and result codes. The value of an EPP result code and HTTP Status-Code MUST remain independent of each other. E.g. an EPP result code indicating an error can be combined with an HTTP request with Status-Code 200 when the EPP command result contains an error, but the HTTP request was successful.

EPP result code: MUST only return EPP result information relating to the EPP protocol. The HTTP header "X-REPP-eppcode" MAY be used to add EPP result information to the HTTP layer.



HTTP Status-Code: MUST only return status information related to the HTTP protocol, When there is a mismatch between the object identifier in the HTTP message-body and the resource URL HTTP Status-Code 412 MUST be returned.

## 9. Command Mapping

This section describes the details of the REST interface by referring to the [\[RFC5730\]](#) Section 2.9 Protocol Commands and defining how these are mapped to RESTful requests.

Each RESTful operation consists of four parts: 1. the resource, 2. the HTTP method 3. the request payload, which is the HTTP message- body of the request, 4. the response payload, being the HTTP message- body of the response.

Table 2 list a mapping for each EPP to REPP, the subsequent sections provide details for each request. All resource URLs in the table are assumed to use the prefix: `"/{context-root}/{version}"/`.

`{c}`: An abbreviation for `{collection}`: this MUST be substituted with "domains", "hosts", "contacts" or "messages".

`{i}`: An abbreviation for `{id}`: this MUST be substituted with the value of a domain name, hostname, contact-id or a message-id.

Command mapping from EPP to REPP.

EPP command	Method	Resource
Hello	GET	/
Login	N/A	N/A
Logout	N/A	N/A
Check	HEAD/GET	/ {c}/ {i}
Info	GET	/ {c}/ {i}
Poll Request	GET	/messages
Poll Ack	DELETE	/messages/ {i}
Transfer Query	GET	/ {c}/ {i}/transfer
Create	POST	/ {c}
Delete	DELETE	/ {c}/ {i}
Renew	PUT	/ {c}/ {i}/validity
Transfer	POST	/ {c}/ {i}/transfer
Transfer Cancel	DELETE	/ {c}/ {i}/transfer

EPP command	Method	Resource
Transfer Approve	PUT	/c/i/transfer
Transfer Reject	DELETE	/c/i/transfer
Update	PUT	/c/i

Table 1

## 9.1. Hello

- Request: GET /{context-root}/{version}
- Request payload: Optional Hello request
- Response payload: Greeting response

The server MUST send a Greeting response, as defined in section 2.4 of [RFC5730] in response to request using the HTTP GET method on the root "/" resource.

A REPP client MAY add content to the HTTP message-body of a Hello request. The version information returned by the server in the Hello response MUST match the version used in the URL of the REPP server.

Example Hello request with an empty message-body:

```
C: GET /repp/v1/ HTTP/1.1
C: Host: repp.example.nl
C: Cache-Control: no-cache
C: Pragma: no-cache
C: Authorization: Bearer <token>
C: Accept: application/epp+xml
C: Accept-Encoding: gzip,deflate
C: Accept-Language: en
C: Accept-Charset: utf-8
c: X-REPP-cltrid: ABC-12345
```

Example Hello response:

```
S: HTTP/1.1 200 OK
S: Date: Fri, 17 Nov 2023 12:00:00 UTC
S: Server: Acme REPP server v1.0
S: Content-Length: 799
S: Content-Type: application/epp+xml
S:
S: <?xml version="1.0" encoding="UTF-8" standalone="no"?>
S: <epp xmlns="urn:ietf:params:xml:ns:epp-1.0">
S:   <greeting>
S:     <!-- rest of the greeting elements -->
S:   </greeting>
S: </epp>
```

## 9.2. Session Management Resources

One of the design goals of REPP is to increase the scalability of an EPP service. This means that session management as described in [RFC5730] is not supported, session management functions are delegated to the HTTP layer.

The server **MUST** not create a client session for use across multiple client requests. The server **MUST** not maintain any state information relating to the client or EPP process state. The client **MUST** include authentication credentials with each request. This **MAY** be done by using any of the available HTTP authentication mechanisms, such as those described in [RFC2617].

### 9.2.1. Login

The Login command defined in [RFC5730] is used to configure a session and is part of the stateful nature of the EPP protocol. A REPP server is stateless and **MUST** not maintain any client state and **MUST** not implement the Login command. The client **MUST** add all the information to a REPP request that is required for the server to be able to properly process the request.

Data attributes from the [RFC5730] Login command are either no longer used or are transferred to the server using a HTTP-header.

- clID: No longer used
- pw: No longer used
- newPW: No longer used
- version: This attribute has been replaced by the version in the server URL.
- lang: This attribute has been replaced by the Accept-Language HTTP request-header.
- svcs: This attribute has been replaced by 1 or more X-REPP-svcs HTTP request-headers.

The server **MUST** check the namespace used in all X-REPP-svcs HTTP request-header. An unsupported namespace **MUST** result in the appropriate EPP result code.

### 9.2.2. Logout

The concept of a session is no longer used by REPP, therefore the Logout command **MUST** not be implemented by the server.

## 9.3. REPP Query Commands

Sending content using an HTTP GET request is discouraged in [RFC9110], there exists no generally defined semantics for content received in a GET request.

A REPP client **MAY** use the HTTP GET method for executing a query command only when no request data has to be added to the HTTP message-body. When an EPP object requires additional authInfo information, as described in [RFC5731] and [RFC5733], the client **MUST** use the HTTP POST method and add the query command content to the HTTP message-body.

### 9.3.1. Check

Request without a request message:

- Request: HEAD {collection}/{id}
- Request payload: N/A
- Response payload: N/A

Request with a request message:

- Request: GET {collection}/{id}
- Request payload: Check request
- Response payload: Check Response

The server MUST support the HTTP HEAD method for the Check command, the server MAY also support the HTTP GET method. If the HTTP HEAD method is used, the client and the server MUST not add any content to the HTTP message-body. If the HTTP GET method is used the client and the server MUST add the Check content to the message-body. The HTTP response for a request using the HTTP GET method, MUST contain the X-REPP-check-avail and MAY contain the X-REPP-check-reason header. The value of X-REPP-check-avail header MUST be "1" or "0" as described in the EPP RFCs, depending on whether the object can be provisioned or not.

A Check request using the HTTP HEAD method is limited to checking only a single resource {id}. This may seem a step backwards when compared to the Check command defined in the EPP RFCs where multiple object-ids are allowed inside a Check command. The RESTful Check command can be load balanced more efficiently when the request contains only a single resource {id} that needs to be checked. When the HTTP GET method is used, the EPP request in the message-body MUST also be limited to a single object to check. The server MUST return EPP result code 2002, when the Check request contains more than 1 object to check.

Example Check request for a domain name:

```
C: HEAD /repp/v1/domains/example.nl HTTP/1.1
C: Host: repp.example.nl
C: Cache-Control: no-cache
C: Pragma: no-cache
C: Authorization: Bearer <token>
C: Accept: application/epp+xml
C: Accept-Encoding: gzip,deflate
C: Accept-Language: en
C: Accept-Charset: utf-8
c: X-REPP-cltrid: ABC-12345
c: X-REPP-svcs: urn:ietf:params:xml:ns:domain-1.0
```

Example Check response:

```
S: HTTP/1.1 200 OK
S: Date: Fri, 17 Nov 2023 12:00:00 UTC
S: Server: Acme REPP server v1.0
S: Content-Length: 0
S: X-REPP-cltrid: ABC-12345
S: X-REPP-svtrid: XYZ-12345
S: X-REPP-check-avail: 0
S: X-REPP-check-reason: In use
```

### 9.3.2. Info

A object <info> request MUST be performed with the HTTP GET or POST method on a resource identifying an object instance. The response MUST be a response message as described in object mapping of the EPP RFCs.

A request for an object without authorization information.

- Request: GET {collection}/{id}
- Request payload: N/A
- Response payload: <info> response

A request for an object that has authorization information attached.

- Request: POST {collection}/{id}
- Request payload: <info> request
- Response payload: <info> response

#### 9.3.2.1. Domain Name

A domain name <info> differs from a contact- and host <info> in the sense that EPP Domain Name Mapping [[RFC5731](#)], Section 3.1.2 describes an OPTIONAL "hosts" attribute for the [domain:name](#) element. This attribute is mapped to additional REST resources to be used in a domain name info request.

The specified default value is "all". This default is mapped to a shortcut, the resource object instance URL without any additional labels.

- default: GET domains/{id}
- Hosts=all: GET domains/{id}/all
- Hosts=del: GET domains/{id}/del
- Hosts=sub: GET domains/{id}/sub
- Hosts=none: GET domains/{id}/none

The server MAY require the client to include additional authorization information. The authorization data MUST be sent with the "X-REPP- authinfo" HTTP request-header.

### 9.3.3. Poll

#### 9.3.3.1. Poll Request

- Request: GET messages/

- Request payload: N/A
- Response payload: Poll request response message.

A client MUST use the HTTP GET method on the messages collection to request the message at the head of the queue.

#### **9.3.3.2. Poll Ack**

- Request: DELETE messages/{id}
- Request payload: N/A
- Response payload: Poll ack response message

A client MUST use the HTTP DELETE method on a message instance to remove the message from the message queue.

#### **9.3.3.3. Transfer Query Op**

- Request: GET {collection}/{id}/transfer
- Request payload: Optional X-REPP-auth-info HTTP header with <authInfo>
- Response payload: Transfer query response message.

A <transfer> query MUST be performed with the HTTP GET method on the transfer resource of a specific object instance.

### **9.4. REPP Transform Commands**

#### **9.4.1. Create**

- Request: POST {collection}/
- Request payload: Object <create>.
- Response payload: Object <create> response.

A client MUST create a new object with the HTTP POST method in combination with an object collection.

#### **9.4.2. Delete**

- Request: DELETE {collection}/{id}
- Request payload: N/A
- Response payload: Object <delete> response.

Deleting an object from the registry database MUST be performed with the HTTP DELETE method on a REST resource specifying a specific object instance.

#### **9.4.3. Renew**

- Request: PUT {collection}/{id}/validity
- Request payload: Object <renew>.
- Response payload: Object <renew> response.

Renewing an object is only specified by [\[RFC5731\]](#), the <renew> command has been mapped to a validity resource.

#### 9.4.4. Update

- Request: PUT {collection}/{id}
- Request payload: Object:update.
- Response payload: Update response message

An object <update> request MUST be performed with the HTTP PUT method on a specific object resource. The payload MUST contain an <object: update> described in the EPP RFCs.

#### 9.4.5. Transfer

Transferring an object from one sponsoring client to another is only specified in [\[RFC5731\]](#) and [\[RFC5733\]](#). The <transfer> command has been mapped to a transfer resource.

The semantics of the HTTP DELETE method are determined by the role of the client executing the method. For the current sponsoring registrar the DELETE method is defined as "reject transfer". For the new sponsoring registrar the DELETE method is defined as "cancel transfer".

##### 9.4.5.1. Create Op

- o Request: POST {collection}/{id}/transfer
- o Request payload: [object:transfer](#).
- o Response Payload: Transfer start response.

Initiating a transfer MUST be done by creating a new "transfer" resource with the HTTP POST method on a specific domain name or contact object instance. The server MAY require authorization information to validate the transfer request.

##### 9.4.5.2. Cancel Op

- Request: DELETE {collection}/{id}/transfer
- Request payload: N/A
- Response payload: Transfer cancel response message.

The new sponsoring client MUST use the HTTP DELETE method to cancel a requested transfer.

##### 9.4.5.3. Approve Op

- Request: PUT {collection}/{id}/transfer
- Request payload: N/A
- Response payload: Transfer approve response message.

The current sponsoring client MUST use the HTTP PUT method to approve a transfer requested by the new sponsoring client.

##### 9.4.5.4. Reject Op

- Request: DELETE {collection}/{id}/transfer

- Request payload: N/A
- Response payload: Transfer reject response message

The current sponsoring client **MUST** use the HTTP DELETE method to reject a transfer requested by the new sponsoring client.

## 10. Transport Considerations

Section 2.1 of the EPP core protocol specification [[RFC5730](#)] describes considerations to be addressed by protocol transport mappings. This document addresses each of the considerations using a combination of features described in this document and features provided by HTTP as follows:

- HTTP is an application layer protocol which uses TCP as a transport protocol. TCP includes features to provide reliability, flow control, ordered delivery, and congestion control. Section 1.5 of [[RFC793](#)] describes these features in detail; congestion control principles are described further in [[RFC2581](#)] and [[RFC2914](#)]. HTTP is a stateless protocol and as such it does not maintain any client state or session.
- The stateful nature of EPP is no longer preserved through managed sessions. There still is a controlled message exchanges because HTTP uses TCP as transport layer protocol.
- HTTP 1.1 allows persistent connections which can be used to send multiple HTTP requests to the server using the same connection. The server **MUST NOT** allow persistent connections.
- The server **MUST NOT** allow pipelining and return EPP result code 2002 if pipelining is detected.
- Batch-oriented processing (combining multiple EPP commands in a single HTTP request) **MUST NOT** be permitted.
- Section 8 of this document describes features to frame EPP request data by adding the data to an HTTP request message-body or request-header.
- A request processing failure has no influence on the processing of other requests. The stateless nature of the server allows a client to retry a failed request or send another request.

## 11. IANA Considerations

TODO: This draft defines three resource collections; domains, contacts, hosts. This may require an IANA RESTful EPP collection protocol registry.

## 12. Internationalization Considerations

TODO

## 13. Security Considerations



[RFC5730] describes a <login> command for transmitting client credentials. This command MUST NOT be used for RESTful EPP. Due to the stateless nature of REST clients MUST transmit their credentials with each HTTP request. The validation of the user credentials must be performed by an out-of-band mechanism. Examples of authentication mechanisms are Basic and Digest access authentication [RFC2617] or OAuth [RFC5849].

To protect data confidentiality and integrity, all data transport between the client and server must use TLS [RFC5246]. Section 9 of [RFC5734] describes the level of security that is REQUIRED.

EPP does not use XML encryption to protect messages. Furthermore, RESTful EPP HTTP servers are vulnerable to common denial-of-service attacks. Therefore, the security considerations of [RFC5734] also apply to RESTful EPP.

## 14. Obsolete EPP Result Codes

The following result codes specified in [RFC5730] are no longer meaningful in RESTful EPP and MUST NOT be used.

Code	Reason
1500	The logout command is not used anymore.
2002	Commands can now be sent in any order.
2100	The REPP URL path includes the version.
2200	The login command is not used anymore.

Table 2

## 15. Acknowledgments

TODO

## 16. Normative References

- [REST] Fielding, R., "Architectural Styles and the Design of Network-based Software Architectures", 2000, <[http://www.ics.uci.edu/~fielding/pubs/dissertation/rest\\_arch\\_style.htm](http://www.ics.uci.edu/~fielding/pubs/dissertation/rest_arch_style.htm)>.
- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <<https://www.rfc-editor.org/info/rfc2119>>.
- [RFC2581] Allman, M., Paxson, V., and W. Stevens, "TCP Congestion Control", RFC 2581, DOI 10.17487/RFC2581, April 1999, <<https://www.rfc-editor.org/info/rfc2581>>.

- 
- [RFC2616]** Fielding, R., Gettys, J., Mogul, J., Frystyk, H., Masinter, L., Leach, P., and T. Berners-Lee, "Hypertext Transfer Protocol -- HTTP/1.1", RFC 2616, DOI 10.17487/RFC2616, June 1999, <<https://www.rfc-editor.org/info/rfc2616>>.
- [RFC2617]** Franks, J., Hallam-Baker, P., Hostetler, J., Lawrence, S., Leach, P., Luotonen, A., and L. Stewart, "HTTP Authentication: Basic and Digest Access Authentication", RFC 2617, DOI 10.17487/RFC2617, June 1999, <<https://www.rfc-editor.org/info/rfc2617>>.
- [RFC2914]** Floyd, S., "Congestion Control Principles", BCP 41, RFC 2914, DOI 10.17487/RFC2914, September 2000, <<https://www.rfc-editor.org/info/rfc2914>>.
- [RFC3375]** Hollenbeck, S., "Generic Registry-Registrar Protocol Requirements", RFC 3375, DOI 10.17487/RFC3375, September 2002, <<https://www.rfc-editor.org/info/rfc3375>>.
- [RFC3986]** Berners-Lee, T., Fielding, R., and L. Masinter, "Uniform Resource Identifier (URI): Generic Syntax", STD 66, RFC 3986, DOI 10.17487/RFC3986, January 2005, <<https://www.rfc-editor.org/info/rfc3986>>.
- [RFC5246]** Dierks, T. and E. Rescorla, "The Transport Layer Security (TLS) Protocol Version 1.2", RFC 5246, DOI 10.17487/RFC5246, August 2008, <<https://www.rfc-editor.org/info/rfc5246>>.
- [RFC5730]** Hollenbeck, S., "Extensible Provisioning Protocol (EPP)", STD 69, RFC 5730, DOI 10.17487/RFC5730, August 2009, <<https://www.rfc-editor.org/info/rfc5730>>.
- [RFC5731]** Hollenbeck, S., "Extensible Provisioning Protocol (EPP) Domain Name Mapping", STD 69, RFC 5731, DOI 10.17487/RFC5731, August 2009, <<https://www.rfc-editor.org/info/rfc5731>>.
- [RFC5732]** Hollenbeck, S., "Extensible Provisioning Protocol (EPP) Host Mapping", STD 69, RFC 5732, DOI 10.17487/RFC5732, August 2009, <<https://www.rfc-editor.org/info/rfc5732>>.
- [RFC5733]** Hollenbeck, S., "Extensible Provisioning Protocol (EPP) Contact Mapping", STD 69, RFC 5733, DOI 10.17487/RFC5733, August 2009, <<https://www.rfc-editor.org/info/rfc5733>>.
- [RFC5734]** Hollenbeck, S., "Extensible Provisioning Protocol (EPP) Transport over TCP", STD 69, RFC 5734, DOI 10.17487/RFC5734, August 2009, <<https://www.rfc-editor.org/info/rfc5734>>.
- [RFC5849]** Hammer-Lahav, E., Ed., "The OAuth 1.0 Protocol", RFC 5849, DOI 10.17487/RFC5849, April 2010, <<https://www.rfc-editor.org/info/rfc5849>>.
- [RFC793]** Postel, J., "Transmission Control Protocol", RFC 793, DOI 10.17487/RFC0793, September 1981, <<https://www.rfc-editor.org/info/rfc793>>.
- [RFC9110]** Fielding, R., Ed., Nottingham, M., Ed., and J. Reschke, Ed., "HTTP Semantics", STD 97, RFC 9110, DOI 10.17487/RFC9110, June 2022, <<https://www.rfc-editor.org/info/rfc9110>>.
-

## Appendix A. Examples

In these examples, lines starting with "C:" represent data sent by a protocol client and lines starting with "S:" represent data returned by a REPP protocol server. Indentation and white space in examples are provided only to illustrate element relationships and are not REQUIRED features of this protocol.

### A.1. X-REPP-auth-info

#### A.1.1. Domain Info with Authorization Data

The X-REPP-auth-info header in a Domain Info Request might look like this: TODO: X-REPP-auth-info must be removed, if auth data is needed, then complete req xml must be sent

```
<?xml version="1.0" encoding="UTF-8" standalone="no"?>
<epp xmlns="urn:ietf:params:xml:ns:epp-1.0">
  <extension>
    <re:rest xmlns:re="urn:ietf:params:xml:ns:restful-epp-1.0">
      <re:authorization>
        <re:pw>passwordfordomain</re:pw>
      </re:authorization>
    </re:rest>
  </extension>
</epp>
```

The HTTP header X-REPP-auth-info MUST contain the entire authorization information element, formatted as described in the EPP RFCs.

### A.2. Domain Create Example

#### A.2.1. Domain Create Request

TODO: remove extension from example below

```
C: POST /repp/v1/domains/ HTTP/1.1
C: Host: repp.example.nl
C: Cache-Control: no-cache
C: Authorization: Basic amRvZTp0ZXN0
C: Pragma: no-cache
C: Accept-Language: en
C: Accept-Charset: utf-8
C: Accept: application/epp+xml
C: X-REPP-cltrid: ABC-12345
c: X-REPP-svcs: urn:ietf:params:xml:ns:domain-1.0
C: Content-Type: text/plain
C: Content-Length: 543

C: <?xml version="1.0" encoding="UTF-8" standalone="no"?>
C: <epp xmlns="urn:ietf:params:xml:ns:epp-1.0"
C:   xmlns:domain="urn:ietf:params:xml:ns:domain-1.0">
C:   <extension>
C:     <re:rest xmlns:re="urn:ietf:params:xml:ns:restful-epp-1.0">
C:       <domain:create>
C:         <!-- Object specific elements-->
C:       </domain:create>
C:     </re:rest>
C:   </extension>
C: </epp>
```

### A.2.2. Domain Create Response:

```
S: HTTP/1.1 200 OK
S: Date: Fri, 17 Nov 2023 12:00:00 UTC
S: Server: Acme REPP server v1.0
S: Content-Language: en
S: Content-Length: 642
S: X-REPP-cltrid: ABC-12345
S: X-REPP-svtrid: 54321-XYZ
S: X-REPP-eppcode: 1000
S: Content-Type: application/epp+xml
S: Connection: close

S: <?xml version="1.0" encoding="UTF-8" standalone="no"?>
S: <epp xmlns="urn:ietf:params:xml:ns:epp-1.0"
S:   xmlns:domain="urn:ietf:params:xml:ns:domain-1.0">
S:   <response>
S:     <result code="1000">
S:       <msg>Command completed successfully</msg>
S:     </result>
S:     <resData>
S:       <domain:creData
S:         <!-- Object specific elements-->
S:       </domain:creData>
S:     </resData>
S:     <trID>
S:       <clTRID>ABC-12345</clTRID>
S:       <svTRID>54321-XYZ</svTRID>
S:     </trID>
S:   </response>
S: </epp>
```

## A.3. Domain Delete Example

### A.3.1. Domain Delete Request:

```
C: DELETE /repp/v1/domains/example.nl HTTP/1.1
C: Host: repp.example.nl
C: Cache-Control: no-cache
C: Authorization: Basic amRvZTp0ZXN0
C: Pragma: no-cache
C: Accept-Language: en
C: Accept-Charset: utf-8
C: X-REPP-cltrid: ABC-12345
```

### A.3.2. Domain Delete Response:

```
S: HTTP/1.1 200 OK
S: Date: Fri, 17 Nov 2023 12:00:00 UTC
S: Server: Acme REPP server v1.0
S: Content-Language: en
S: Content-Length: 505
S: X-REPP-cltrid: ABC-12345
S: X-REPP-svtrid: 54321-XYZ
S: X-REPP-eppcode: 1000
S: Content-Type: application/epp+xml
S: Connection: close

S: <?xml version="1.0" encoding="UTF-8" standalone="no"?>
S: <epp xmlns="urn:ietf:params:xml:ns:epp-1.0"
S:   xmlns:domain="urn:ietf:params:xml:ns:domain-1.0">
S:   <response>
S:     <result code="1000">
S:       <msg>Command completed successfully</msg>
S:     </result>
S:     <trID>
S:       <clTRID>ABC-12345</clTRID>
S:       <svTRID>54321-XYZ</svTRID>
S:     </trID>
S:   </response>
S: </epp>
```

## Authors' Addresses

### Maarten Wullink

SIDN Labs

Email: [maarten.wullink@sidn.nl](mailto:maarten.wullink@sidn.nl)

URI: <https://sidn.nl/>

### Marco Davids

SIDN Labs

Email: [marco.davids@sidn.nl](mailto:marco.davids@sidn.nl)

URI: <https://sidn.nl/>